

Stat 33B - Lecture 6

February 26, 2020

Announcements

Quiz 2 this Friday:

- Similar length to Quiz 1
- Multiple-choice and fill-in-the-blank questions.
- Covers everything up to the end of last week.

Review

Previous lecture:

- Tidy data
- The `pivot_wider` function

We took this untidy data:

```
library(tidyr)
```

```
table2
```

```
## # A tibble: 12 x 4
##   country      year type      count
##   <chr>      <int> <chr>    <int>
## 1 Afghanistan 1999 cases      745
## 2 Afghanistan 1999 population 19987071
## 3 Afghanistan 2000 cases      2666
## 4 Afghanistan 2000 population 20595360
## 5 Brazil      1999 cases      37737
## 6 Brazil      1999 population 172006362
## 7 Brazil      2000 cases      80488
## 8 Brazil      2000 population 174504898
## 9 China       1999 cases      212258
## 10 China       1999 population 1272915272
## 11 China       2000 cases      213766
## 12 China       2000 population 1280428583
```

And used the `pivot_wider()` function to rotate rows into columns:

```
pivot_wider(table2, names_from = type, values_from = count)
```

```
## # A tibble: 6 x 4
##   country      year cases population
##   <chr>      <int> <int>    <int>
## 1 Afghanistan 1999     745  19987071
## 2 Afghanistan 2000    2666  20595360
## 3 Brazil      1999   37737  172006362
## 4 Brazil      2000   80488  174504898
```

```
## 5 China      1999 212258 1272915272
## 6 China      2000 213766 1280428583
```

Rotating Columns into Rows

Some data sets put measurements for multiple observations in a single row.

In this table, each row contains the measurements for two observations of the `cases` variable:

```
table4a
```

```
## # A tibble: 3 x 3
##   country    `1999` `2000`
## * <chr>      <int> <int>
## 1 Afghanistan    745   2666
## 2 Brazil         37737  80488
## 3 China          212258 213766
```

To make this data tidy:

- Rotate the 1999 and 2000 columns into rows.
- New columns are `year` and `cases`.

Tidying this data set makes it longer.

Important Note: Earlier we saw negative indexing with positions:

```
table4a[-1]
```

```
## # A tibble: 3 x 2
##   `1999` `2000`
##   <int> <int>
## 1     745   2666
## 2   37737  80488
## 3 212258 213766
```

Tidverse functions (but not base R!) also support negative indexing with names!

Use `pivot_longer()` to rotate columns into rows:

```
pivot_longer(table4a, -country, names_to = "year", values_to = "cases")
```

```
## # A tibble: 6 x 3
##   country    year  cases
##   <chr>      <chr> <int>
## 1 Afghanistan 1999     745
## 2 Afghanistan 2000    2666
## 3 Brazil      1999   37737
## 4 Brazil      2000   80488
## 5 China       1999  212258
## 6 China       2000  213766
```

Need to specify:

- Columns to rotate as `cols`.
- Name(s) of new identifier column(s) as `names_to`.
- Name(s) of new measurement column(s) as `values_to`.

See `vignette("pivot")` for more examples of using `tidyr`.

If you wanted to do this without `tidyr`:

1. Subset columns to separate 1999 and 2000 into two data frames.

2. Add a `year` column to each.
3. Rename the 1999 and 2000 columns to `cases`.
4. Stack the two data frames with `rbind()`.

```
col199 = table4a[-3]
col100 = table4a[-2]

names(col199)[2] = "cases"
col199["year"] = 1999

col199
```

```
## # A tibble: 3 x 3
##   country      cases  year
##   <chr>      <int> <dbl>
## 1 Afghanistan    745  1999
## 2 Brazil       37737  1999
## 3 China       212258  1999
```

```
names(col100)[2] = "cases"
col100["year"] = 2000

col100
```

```
## # A tibble: 3 x 3
##   country      cases  year
##   <chr>      <int> <dbl>
## 1 Afghanistan   2666  2000
## 2 Brazil       80488  2000
## 3 China       213766  2000
```

```
rbind(col199, col100)
```

```
## # A tibble: 6 x 3
##   country      cases  year
##   <chr>      <int> <dbl>
## 1 Afghanistan    745  1999
## 2 Brazil       37737  1999
## 3 China       212258  1999
## 4 Afghanistan   2666  2000
## 5 Brazil       80488  2000
## 6 China       213766  2000
```

Merging Data

A merge or “join” combines data from two separate data frames, based on some identifying values they have in common.

Recall the `pivot_wider()` example with `table2`.

To make `table2` tidy without `tidyr`:

1. Subset rows to separate `cases` and `population` values.
2. Remove the `type` column from each.
3. Rename the `count` column to `cases` and `population`.
4. Merge these two subsets by matching `country` and `year`.

The two subsets are:

```
table2
```

```
## # A tibble: 12 x 4
##   country      year type      count
##   <chr>      <int> <chr>    <int>
## 1 Afghanistan 1999 cases      745
## 2 Afghanistan 1999 population 19987071
## 3 Afghanistan 2000 cases      2666
## 4 Afghanistan 2000 population 20595360
## 5 Brazil      1999 cases      37737
## 6 Brazil      1999 population 172006362
## 7 Brazil      2000 cases      80488
## 8 Brazil      2000 population 174504898
## 9 China       1999 cases      212258
## 10 China      1999 population 1272915272
## 11 China      2000 cases      213766
## 12 China      2000 population 1280428583
```

```
cases = table2[table2$type == "cases", ]
pop = table2[table2$type == "population", ]
```

```
names(pop)[4] = "population"
pop = pop[-3]
```

```
pop
```

```
## # A tibble: 6 x 3
##   country      year population
##   <chr>      <int>    <int>
## 1 Afghanistan 1999    19987071
## 2 Afghanistan 2000    20595360
## 3 Brazil      1999    172006362
## 4 Brazil      2000    174504898
## 5 China       1999   1272915272
## 6 China       2000   1280428583
```

```
names(cases)[4] = "cases"
cases = cases[-3]
```

BAD ways to combine rows:

```
cases$population = pop$population
```

```
# OR
```

```
cbind(cases, pop["population"])
```

Dangerous if you're not 110% sure the order of rows matches!

Imagine if the rows were in a different order:

```
pop = pop[c(4, 5, 1, 2, 3, 6), ]
```

Then adding a new column or using `cbind()` would mix up the observations.

Instead, match rows on `country` and `year` with the `merge()` function:

```
merge(cases, pop)
```

```
##      country year  cases population
## 1 Afghanistan 1999    745   19987071
## 2 Afghanistan 2000   2666   20595360
## 3      Brazil 1999  37737  172006362
## 4      Brazil 2000  80488  174504898
## 5      China 1999 212258 1272915272
## 6      China 2000 213766 1280428583
```

The function automatically uses columns that have the same name in both tables to match.

More Merges

Data split across multiple tables are called *relational data*.

A column shared by several tables is called a *key*.

For example, a grocery store's inventory system might have:

- A table that lists stores
- A table that lists items (fruits, vegetables, etc)
- A table that lists quantity of each item at each store

See the bCourse for the data set:

```
stores = readRDS("data/grocery/stores.rds")
```

```
stores
```

```
##  store_id  manager status      city
## 1         1      Chen    20    Oakland
## 2         2 Hernandez    10 San Francisco
## 3         3     Smith    30 San Francisco
## 4         4       Ali    20    Oakland
## 5         5     Rossi    30    Berkeley
```

```
items = readRDS("data/grocery/items.rds")
```

```
items
```

```
##  ItemID ItemName      Variety Price  Source
## 1      1   Lemon      Eureka  0.50 San Diego
## 2      2  Orange    Valencia  1.00 San Diego
## 3      3  Apple Red Delicious  0.90  Spokane
## 4      4  Apple      Fuji    0.90  Spokane
## 5      5   Lime      Key    1.50 Veracruz
## 6      6  Mango    Alfonso  2.25  Lucknow
```

```
inv = readRDS("data/grocery/inventory.rds")
```

```
inv
```

```
##  StoreID ItemID Qty
## 1      1      1 300
## 2      1      2 200
## 3      1      3 400
## 4      1      4 200
## 5      1      5 100
```

```
## 6      1      6 100
## 7      2      1 300
## 8      2      2 400
## 9      3      2 200
## 10     4      2 200
## 11     4      4 300
## 12     4      5 400
```

By default, the `merge()` function only keeps rows that match:

```
merge(stores, inv)
```

```
##      store_id  manager status      city StoreID ItemID Qty
## 1          1      Chen    20    Oakland        1      1 300
## 2          2 Hernandez    10 San Francisco        1      1 300
## 3          3      Smith    30 San Francisco        1      1 300
## 4          4        Ali    20    Oakland        1      1 300
## 5          5      Rossi    30    Berkeley        1      1 300
## 6          1      Chen    20    Oakland        1      2 200
## 7          2 Hernandez    10 San Francisco        1      2 200
## 8          3      Smith    30 San Francisco        1      2 200
## 9          4        Ali    20    Oakland        1      2 200
## 10         5      Rossi    30    Berkeley        1      2 200
## 11         1      Chen    20    Oakland        1      3 400
## 12         2 Hernandez    10 San Francisco        1      3 400
## 13         3      Smith    30 San Francisco        1      3 400
## 14         4        Ali    20    Oakland        1      3 400
## 15         5      Rossi    30    Berkeley        1      3 400
## 16         1      Chen    20    Oakland        1      4 200
## 17         2 Hernandez    10 San Francisco        1      4 200
## 18         3      Smith    30 San Francisco        1      4 200
## 19         4        Ali    20    Oakland        1      4 200
## 20         5      Rossi    30    Berkeley        1      4 200
## 21         1      Chen    20    Oakland        1      5 100
## 22         2 Hernandez    10 San Francisco        1      5 100
## 23         3      Smith    30 San Francisco        1      5 100
## 24         4        Ali    20    Oakland        1      5 100
## 25         5      Rossi    30    Berkeley        1      5 100
## 26         1      Chen    20    Oakland        1      6 100
## 27         2 Hernandez    10 San Francisco        1      6 100
## 28         3      Smith    30 San Francisco        1      6 100
## 29         4        Ali    20    Oakland        1      6 100
## 30         5      Rossi    30    Berkeley        1      6 100
## 31         1      Chen    20    Oakland        2      1 300
## 32         2 Hernandez    10 San Francisco        2      1 300
## 33         3      Smith    30 San Francisco        2      1 300
## 34         4        Ali    20    Oakland        2      1 300
## 35         5      Rossi    30    Berkeley        2      1 300
## 36         1      Chen    20    Oakland        2      2 400
## 37         2 Hernandez    10 San Francisco        2      2 400
## 38         3      Smith    30 San Francisco        2      2 400
## 39         4        Ali    20    Oakland        2      2 400
## 40         5      Rossi    30    Berkeley        2      2 400
## 41         1      Chen    20    Oakland        3      2 200
## 42         2 Hernandez    10 San Francisco        3      2 200
```

## 43	3	Smith	30	San Francisco	3	2	200
## 44	4	Ali	20	Oakland	3	2	200
## 45	5	Rossi	30	Berkeley	3	2	200
## 46	1	Chen	20	Oakland	4	2	200
## 47	2	Hernandez	10	San Francisco	4	2	200
## 48	3	Smith	30	San Francisco	4	2	200
## 49	4	Ali	20	Oakland	4	2	200
## 50	5	Rossi	30	Berkeley	4	2	200
## 51	1	Chen	20	Oakland	4	4	300
## 52	2	Hernandez	10	San Francisco	4	4	300
## 53	3	Smith	30	San Francisco	4	4	300
## 54	4	Ali	20	Oakland	4	4	300
## 55	5	Rossi	30	Berkeley	4	4	300
## 56	1	Chen	20	Oakland	4	5	400
## 57	2	Hernandez	10	San Francisco	4	5	400
## 58	3	Smith	30	San Francisco	4	5	400
## 59	4	Ali	20	Oakland	4	5	400
## 60	5	Rossi	30	Berkeley	4	5	400

Use the `all` parameter to force all rows from both tables to show up:

```
merge(stores, inv, by.x = "store_id", by.y = "StoreID", all = TRUE)
```

##	store_id	manager	status	city	ItemID	Qty
## 1	1	Chen	20	Oakland	1	300
## 2	1	Chen	20	Oakland	2	200
## 3	1	Chen	20	Oakland	3	400
## 4	1	Chen	20	Oakland	4	200
## 5	1	Chen	20	Oakland	5	100
## 6	1	Chen	20	Oakland	6	100
## 7	2	Hernandez	10	San Francisco	1	300
## 8	2	Hernandez	10	San Francisco	2	400
## 9	3	Smith	30	San Francisco	2	200
## 10	4	Ali	20	Oakland	2	200
## 11	4	Ali	20	Oakland	4	300
## 12	4	Ali	20	Oakland	5	400
## 13	5	Rossi	30	Berkeley	NA	NA

There are also `all.x` and `all.y` parameters to only force all rows from one table.

Use the `by` parameter to specify the key. If the name of the key is different for the two tables, use `by.x` and `by.y`.

```
merge(stores, inv, by.x = "store_id", by.y = "StoreID")
```

##	store_id	manager	status	city	ItemID	Qty
## 1	1	Chen	20	Oakland	1	300
## 2	1	Chen	20	Oakland	2	200
## 3	1	Chen	20	Oakland	3	400
## 4	1	Chen	20	Oakland	4	200
## 5	1	Chen	20	Oakland	5	100
## 6	1	Chen	20	Oakland	6	100
## 7	2	Hernandez	10	San Francisco	1	300
## 8	2	Hernandez	10	San Francisco	2	400
## 9	3	Smith	30	San Francisco	2	200
## 10	4	Ali	20	Oakland	2	200
## 11	4	Ali	20	Oakland	4	300

```
## 12      4      Ali      20      Oakland      5 400
```

If you're familiar with SQL JOINS, the `merge()` function is the same idea. Specifically:

- INNER JOIN is `all = FALSE` (the default)
- LEFT JOIN is `all.x = TRUE`
- RIGHT JOIN is `all.y = TRUE`
- OUTER JOIN is `all = TRUE`